

Client Demo Recording Guide

Purpose: Self-record the UI walkthrough that goes into the client deck. Target length **4–5 minutes**.

This is the centerpiece of Section 2 in [Argos_VSP_Client_52slides_Apr2026.pptx](#). When you finish, save the source `.mp4` at:

```
presentation_materials_20260224/06_demo_videos/_ui_walkthrough_clientpitch.mp4
```

1. Equipment & Setup

Recorder

- **Loom** (<https://www.loom.com/>) — free tier is enough; auto-uploads, click highlighting, easy trim. **Recommended.**
- **macOS QuickTime Player** — File → New Screen Recording. Built-in, no install.
- **OBS Studio** — most control, free. Install the "Highlight Cursor" plugin separately.
- **Avoid Zoom recordings** — quality is too low for an embedded slide.

Microphone

- USB headset mic or external USB mic preferred (Blue Yeti / Snowball / similar).
- The MacBook built-in mic is acceptable in a quiet room — but **do a 30-second test recording first** and listen back with headphones. If you hear any room echo or fan noise, switch to a headset mic.

Camera (optional)

- A picture-in-picture of your face in a corner can help, but it's not required and can be distracting.
- Default to no camera unless you're confident on camera.

Browser

- Fresh **Chrome** window — no extensions visible.
- Full-screen.
- **Bookmarks bar hidden** (`Cmd+Shift+B` on macOS to toggle).
- No other tabs open.

Resolution

- **1920×1080 minimum.** In Loom, set "HD 1080p" before starting.
- In QuickTime, click the down-arrow next to the record button → make sure microphone is selected.

Cursor highlighting

- **On.** Loom does this automatically.
- OBS: install the "Highlight Cursor" plugin.
- QuickTime: shows a click ring by default if "Show Mouse Clicks" is enabled in the recording dropdown.

2. Pre-Record Checklist (10 min)

Don't skip these. The most expensive thing in this 48 hours is a recording where the demo crashed.

1. **Pre-run the pipeline once** end-to-end on the same input video you're about to record with. Confirm nothing crashes. Confirm the report opens. Confirm the tier badges (Trust / Salvage / Strip) and per-word coloring (blue / orange / purple) both render.
2. **Stage the input video** — pick one with clear speech and one or two named entities. Confirm it produces a high-IS, mostly-blue output with at least one Trust-tier segment, plus ideally one Salvage and one Strip segment elsewhere in the report so you can demo all three tiers in Scene 3.
3. **Do not gamble on a fresh clip during the recording.**
3. **Decode with `VSP_NBEST=1`** so the joint conf + beam-agreement rule and tier badges actually fire. Without it, the report falls back to a conf-only coloring with no tier pills — which is *not* the story we're telling clients.
4. **Clear the output folder** so the Complete screen looks fresh during the demo.
5. **Close Slack, email, calendar, notifications.** macOS Do Not Disturb on. Quit anything that might pop up.
6. **Set browser zoom 110–125%** so text is readable in the recorded video.
7. **One full dry-run aloud, with a timer.** Aim under 5 minutes. If you go over, cut something.

3. The Script — 4 Scenes, ~1 Min Each

Read these out loud while you click. **Pause 1–2 seconds between scenes** to make editing easy later.

Scene 1 — The setup (0:00–0:45)

"This is the visual speech recognition pipeline. The whole product runs in your browser. I'm going to take a video, drop it in, and walk you through what comes out the other end. Nothing's pre-cached — this is the real thing running live."

(Drag-drop your test video onto the page. Let the validation animation play.)

"The system has detected the video, validated it, and is showing me how many segments it'll process."

Scene 2 — The pipeline running (0:45–2:00)

(Click "Start Processing." Let the progress bar run.)

"It's now running through nine automatic stages — face detection, mouth cropping, ASR, LLM-based decoding. You can watch progress in real time. For your content this would take roughly N minutes per hour of video."

(Edit out long waits when you cut. You don't need to show all 9 stages live; cut to "almost done" and show the final stages.)

Scene 3 — The report (2:00–3:30) — *most important scene*

(Click "Open Results." Open the HTML report.)

"Here's the part most clients care about. Every segment tells you twice how much to trust it — once at the **segment level** with a tier badge, and once at the **word level** with per-word coloring."

(Point to the rightmost column.)

"Each segment is tagged **Trust**, **Salvage**, or **Strip**. Trust means the segment is solidly transcribed — read it normally. Salvage means it's useful but partial — verify names and numbers. Strip means we don't trust it enough to color the words at all, so we grey them out — skim for gist, don't quote."

(Scroll to a Trust segment. Hover over a blue word, then an orange word, then a purple word.)

"Inside a segment, words are **blue** for high confidence, **orange** for medium, **purple** for low. The blue threshold isn't just probability — the model also has to agree with itself across multiple decoding hypotheses. That second signal catches confident-but-wrong fluent guesses, which is the most dangerous failure mode in lip reading."

(Scroll to a Strip-tier segment to show the grey-italic stripped coloring + banner.)

"When the segment-level confidence falls below our reliability floor, we strip the coloring entirely — because at that point the per-word signal is misleading. The system tells you *not* to trust it, instead of pretending it's sure."

(Point to the IS column.)

"On the right is the segment-level Intelligibility Score. Calibrated against expert judgment — above 3.8 is clearly conveyed, between 2 and 3.8 is useful with context, below 2 we flag for review."

Scene 4 — The trust message (3:30–4:30)

(Scroll to a Strip-tier or hallucinated segment.)

"Here's a segment the system flagged as unreliable. The model knew it wasn't sure — Strip badge, low IS, coloring removed, automatically pulled out of the 'usable' bucket. **You don't have to review every segment to find the bad ones; the system finds them for you.** That's the difference between this and a black-box transcript."

(Optionally point to a number or a proper name in any segment.)

"One special rule worth knowing: numbers, names, dates, and dollar amounts are always capped at orange — even when the model is sure. Lip reading physically cannot disambiguate 'fifteen' from 'fifty' or 'million' from 'billion'. The system refuses to mark those blue, so you always know to verify them against the source."

(Pause. End on the dashboard or Complete screen.)

"That's the whole pipeline. Drag in, walk away, come back to color-coded results you can trust — with the system telling you exactly which words and segments to trust."

4. Common Pitfalls

- **Don't apologize** for anything that's working. If a stage takes 30 seconds, say "this stage takes 30 seconds" — don't say "sorry this is slow." Confidence is contagious.
- **Don't read out the technical names** ("VSP-LLM decoder," "k-means clustering") to a non-technical client. Say what those things *do*.
- **Cut, don't restart.** If you flub a sentence, pause for 2 seconds, then say it again cleanly. You'll edit out the mistake.
- **Speak ~20% slower than feels natural.** Watch a previous recording — almost everyone speaks too fast on the first take.
- **Hover, don't just click.** Move your cursor *to* a word *before* clicking. The viewer's eye follows your cursor; help them.
- **End on a still frame**, not mid-click. The last thing the client sees should be a clean Complete or report screen.

5. Editing & Embedding

- **Trim** dead air at the start/end and between scenes. Loom has a built-in trim. iMovie or DaVinci Resolve (free) for anything more complex.
- **Export as MP4 / H.264 / 1080p.**
- **Embed in PowerPoint with Insert → Video → Video on My PC**, *not* a link — clients may not have internet during the meeting.

- **Save the source .mp4** at `presentation_materials_20260224/06_demo_videos/_ui_walkthrough_clientpitch.mp4` so it's tracked alongside the example clips.
-

6. Quality Checklist Before Embedding

- Audio levels consistent across scenes (no sudden volume jumps).
- Cursor visible and highlighted throughout.
- No notification pop-ups visible in the recording.
- Total length under 5 minutes.
- Plays at 1080p when previewed full-screen.
- Embeds (not links) into the PPTX — plays offline.