

# Q&A; Cheat Sheet — One Page

Closing line — say this if nothing else lands:

*"We are not asking you to trust the model. We give your reviewers the model's output, its uncertainty, its failure flags, and the evidence needed to decide what deserves human attention."*

## Three anchor numbers (the headline)

**62% review-useful · 23% clearly conveyed · 1 in 5 auto-flagged.** Full 1,497-segment unfiltered real-world baseline. Never say "62% accurate" — say "62% review-useful."

## Three lines to land

1. **"Confidence is triage, not truth."** Green = review faster. Red = review first.
2. **"A wrong fluent transcript is worse than an empty one. Our value is not that failure never happens — it's that failure is surfaced."**
3. **"The model is a high-quality first draft with calibrated uncertainty. The human triages."**

## Top Q&A;

| Q                                           | A                                                                                                                                                                                     |
|---------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| How accurate?                               | "Depends on clip quality. On 1,497 real-world segments, ~62% review-useful, ~23% clearly conveyed. I won't promise a number on your footage before we run your clips."                |
| Can we trust the transcript?                | "Not blindly — that's the point. Word + segment confidence + hallucination flags. Trust high-confidence faster, review yellow/red, exclude low-confidence from auto-conclusions."     |
| What about hallucinations?                  | "Core failure mode, not edge case. Token confidence + length anomalies flag them. The Obama 'Rwanda's genocide' case: caught at prob 0.02. <i>Surfaced, not hidden.</i> "             |
| Was the 'reviewer' a human?                 | "No — blind frontier-LLM evaluator at scale. Right next step is human validation on your pilot clips. Say 'blind LLM evaluator,' never 'expert reviewer.'"                            |
| Two people talking?                         | "Today: one centered speaker. Path: speaker-specific cropping, decode each independently. Engineering, not research. ~3 weeks to ablation on your data."                              |
| Evidence use?                               | "Only with human review and confidence metadata. Evidence-discovery and review-acceleration tool — not standalone evidence."                                                          |
| Cherry-picked examples?                     | "Curated for teaching range; the headline numbers are full 1,497-segment, not selected. We can run your clips next."                                                                  |
| Is green always correct?                    | "No. 92.8% reliable in high-quality segments, 21.8% in low. UI policy: ≥0.82 full coloring, 0.65–0.82 amber banner, <0.65 strip. <i>The UI removes coloring where it would lie.</i> " |
| What's different vs other lip-reading work? | "Real-world end-to-end pipeline (not LRS3-only); uncertainty surfaced at every level; deployable infrastructure cloud-or-on-prem — not a research codebase."                          |
| Human IS comparison?                        | "Lay ~0.9, deaf adult ~2.7, forensic expert ~3.0, model alone 2.52, <b>reviewer + model ~3.8</b> . Alone we tie a deaf adult; combined system beats the expert."                      |
| vs 5-expert team?                           | "Same quality (~3.8 IS), <b>50–100× cheaper</b> (\$3–6/min vs \$200–400/min), structurally more reproducible. One reviewer + Argos, not five experts."                                |
| Could it be misused?                        | "Yes — that's why we ship it as reviewable evidence stream, not autonomous truth. Won't deploy without human verification. Deal-breaker for us."                                      |
| What's a pilot?                             | "3–5 canonical clips from your domain, two weeks, full pipeline + confidence-scored report, your reviewer evaluates fit. No infra ask up front."                                      |

## Don't say → say instead

| Don't say                          | Say instead                                        |
|------------------------------------|----------------------------------------------------|
| "62% accurate"                     | "62% review-useful"                                |
| "Expert reviewer"                  | "blind LLM evaluator"                              |
| "Near-zero hallucination risk"     | "hallucinations flagged before review"             |
| "The system knows when it's right" | "it tells you where to look first"                 |
| "Replaces human review"            | "routes review to the right places"                |
| "Replaces 5 experts"               | "replaces 5 experts with one reviewer + the model" |
| "State of the art"                 | don't claim it                                     |
| "Model beats human lip readers"    | "combined system beats unaided humans"             |

## Pull-only (don't volunteer)

**Beam aggregation:** shipped May 1, env-var gated, full 1,497 eval running. WER 59→57% on 107-segment validation; voting method has best calibration error of any method measured.

**Runtime confidence:** per-token softmax + length-anomaly check. **IS is dev-time only — needs reference text. Never claim IS at runtime.**